

Homework 15

- Exercise B.21 in Christensen.
- Let \mathbf{X} denote the $n \times p$ design matrix where $\mathbf{X}^T \mathbf{1} = \mathbf{0}$ and partition $\mathbf{X} = [\mathbf{X}_\gamma \mathbf{X}_{1-\gamma}]$ where \mathbf{X}_γ denotes the design matrix under model γ which includes the columns of \mathbf{X} where $\gamma_j = 1$ (and intercept) $\mathbf{X}_{1-\gamma}$ are the columns of \mathbf{X} where $\gamma_j = 0$. Similarly partition $\boldsymbol{\beta} = (\boldsymbol{\beta}_\gamma^T, \boldsymbol{\beta}_{1-\gamma}^T)^T$ the vector of regression coefficients. Using Zellner's g prior for

$$\boldsymbol{\beta} \mid g, \sigma^2 \sim \mathbf{N}(\mathbf{0}, g\sigma^2(\mathbf{X}^T \mathbf{X})^{-1})$$

find $\boldsymbol{\beta}_\gamma \mid g, \sigma^2, \boldsymbol{\beta}_{1-\gamma} = \mathbf{0}$ and show that this is also a g prior. *Hint use Problem 1*

- Refer to the Womack et al (2014) paper <http://amstat.tandfonline.com/doi/abs/10.1080/01621459.2014.880348>). Show that the Bayes factor (equation 14) for comparing \mathcal{M}_γ to the null model may be written in terms of the Appell F_1 function:

$$F_1(a, b_1, b_2, c; x, y) = \frac{\Gamma(c)}{\Gamma(a)\Gamma(c-a)} \int_0^1 t^{a-1} (1-t)^{c-a-1} (1-xt)^{-b_1} (1-yt)^{-b_2} dt, \quad c > a > 0$$

(see <http://functions.wolfram.com/HypergeometricFunctions/AppellF1/07/01/01/0001/> where the prior on w is a Beta(1/2, 1/2) Similarly find an expression for the posterior mean of $g/(1+g)$, the posterior expected shrinkage.

- For extra participation credit: Use the **Appell** package and write two functions to return the BF and the posterior expected shrinkage respectively using R2, number of predictors and n as input. Compare results to Womack et al. Post on Piazza.
- Refer back to the simulation study you did using the Nott-Kohn code in HW 11. Use the `prior="hyper-g-n"` prior with `a = 3` and the Zellner-Siow Cauchy priors `prior="ZS-null", a=n` where n is the sample size in the **BAS** package to estimate $\boldsymbol{\beta}$ under enumeration of all models BMA (see `help(coef.bma)`), the Robust prior of Bayarri et al (R package **BayesVarSel** or **BVS** and the (optionally) intrinsic prior (use your code or the archived R package **VarselIP**) How do these compare to your previous results in terms of bias or MSE and on running time? **BAS** should be able to enumerate all models `n.models=2^p` where p = number of predictors.